

Code challenge: solutions

Nathan Green

nathan.green@imperial.ac.uk

These are some example model solutions to the code challenge, set as part of the R for trial & model-based cost-effectiveness analysis workshop taking place 9th July 2019 at University College London. The questions can be found [here](#).

Load supporting packages.

```
library(purrr)
library(knitr)
```

C1. A simple decision tree

Define the separate monthly transition probabilities.

```
delta0 <- 0.001182
deltac <- 0.025
deltaa <- 0.08
betac <- 0.0027
betaa <- 0.0083
```

and the utilities of being in each state.

```
sutils <- c(dead = 0, CA = 0.3, cancer = 0.6, AIDS = 0.5, well = 1)
```

Note that the order of states is dead, cancer & AIDS, cancer, AIDS, well.

Define decision trees in terms of the structure, values and probabilities. We'll use the `tribble` function just because it allows us to specify a matrix by rows rather than columns.

```
library(tibble)

tree_probs <- list()

# unique state outcomes each branch
tree_probs$well <-
  tribble(~rowname, ~dead, ~ndead, ~recr, ~ncancer, ~CA, ~cancer, ~AIDS, ~well,
    "well0", delta0, 1-delta0, NA, NA, NA, NA, NA, NA,
    "ndead", NA, NA, betac, 1-betac, NA, NA, NA, NA,
    "recr", NA, NA, NA, NA, betaa, 1-betaa, NA, NA,
    "ncancer", NA, NA, NA, NA, NA, NA, betaa, 1-betaa) %>%
  column_to_rownames()

tree_probs$cancer <-
  tribble(~rowname, ~dead, ~ndead, ~diec, ~survc, ~CA, ~cancer,
    "cancer0", delta0, 1-delta0, NA, NA, NA, NA,
    "ndead", NA, NA, deltac, 1-deltac, NA, NA,
    "survc", NA, NA, NA, NA, betaa, 1-betaa) %>%
  column_to_rownames()

tree_probs$AIDS <-
  tribble(~rowname, ~dead, ~ndead, ~diea, ~surva, ~CA, ~AIDS,
    "AIDS0", delta0, 1-delta0, NA, NA, NA, NA,
    "ndead", NA, NA, deltaa, 1-deltaa, NA, NA,
```

```

      "surva", NA,      NA,      NA,      NA,      betac,1-betac) %>%
column_to_rownames()

tree_probs$CA <-
tribble(~rowname, ~dead, ~ndead, ~diec, ~survc, ~diea, ~CA,
        "CA0",    delta0,1-delta0,NA,      NA,      NA,      NA,
        "ndead",  NA,      NA,      deltac,1-deltac,NA,      NA,
        "survc",  NA,      NA,      NA,      NA,      deltaa,1-deltaa) %>%
column_to_rownames()

tree_probs$dead <- 1

tree_probs

## $well
##           dead      ndead recurc ncancer      CA cancer      AIDS      well
## well0  0.001182 0.998818      NA      NA      NA      NA      NA      NA
## ndead      NA      NA 0.0027 0.9973      NA      NA      NA      NA
## recurc      NA      NA      NA      NA 0.0083 0.9917      NA      NA
## ncancer      NA      NA      NA      NA      NA      NA 0.0083 0.9917
##
## $cancer
##           dead      ndead diec survc      CA cancer
## cancer0 0.001182 0.998818      NA      NA      NA      NA
## ndead      NA      NA 0.025 0.975      NA      NA
## survc      NA      NA      NA      NA 0.0083 0.9917
##
## $AIDS
##           dead      ndead diea surva      CA      AIDS
## AIDS0 0.001182 0.998818      NA      NA      NA      NA
## ndead      NA      NA 0.08 0.92      NA      NA
## surva      NA      NA      NA      NA 0.0027 0.9973
##
## $CA
##           dead      ndead diec survc diea      CA
## CA0 0.001182 0.998818      NA      NA      NA      NA
## ndead      NA      NA 0.025 0.975      NA      NA
## survc      NA      NA      NA      NA 0.08 0.92
##
## $dead
## [1] 1

tree_vals <- list()

tree_vals$well <-
tribble(~rowname, ~dead, ~ndead, ~cancer, ~ncancer, ~CA, ~CnotA, ~AIDS, ~well,
        "well0", 0,0,0,0,0,0,0,0,0,
        "ndead", 0,0,0,0,0,0,0,0,0,
        "cancer", 0,0,0,0,0.3,0.6,0,0,
        "ncancer", 0,0,0,0,0,0,0.5,1) %>%
column_to_rownames()

tree_vals$cancer <-
tribble(~rowname, ~dead, ~ndead, ~cancer, ~ncancer, ~CA, ~CnotA, ~AIDS, ~well,

```

```

      "well", 0,0,0,0,0,0,0,0,
      "ndead", 0,0,0,0,0,0,0,0,
      "cancer", 0,0,0,0,0.3,0.6,0,0,
      "ncancer", 0,0,0,0,0,0,0.5,1) %>%
column_to_rownames()

tree_vals$AIDS <-
  tribble(~rowname, ~dead, ~ndead, ~cancer, ~ncancer, ~CA, ~CnotA, ~AIDS, ~well,
    "well", 0,0,0,0,0,0,0,0,
    "ndead", 0,0,0,0,0,0,0,0,
    "cancer", 0,0,0,0,0.3,0.6,0,0,
    "ncancer", 0,0,0,0,0,0,0.5,1) %>%
column_to_rownames()

tree_vals$CA <-
  tribble(~rowname, ~dead, ~ndead, ~cancer, ~ncancer, ~CA, ~CnotA, ~AIDS, ~well,
    "well", 0,0,0,0,0,0,0,0,
    "ndead", 0,0,0,0,0,0,0,0,
    "cancer", 0,0,0,0,0.3,0.6,0,0,
    "ncancer", 0,0,0,0,0,0,0.5,1) %>%
column_to_rownames()

```

C2. Extend C1 for multiple cycles

Assuming a binomial tree we can forward simulate for a synthetic cohort. This is a brute force approach and is potentially time-consuming.

```

cohort <- list()
n_cohort <- 1000
death_states <- c("diea", "diec", "dead")

for (i in seq_len(n_cohort)) {

  traj_s <- NULL
  traj_u <- NULL
  state_name <- "well"

  while (!state_name %in% death_states) {

    p <- tree_probs[[state_name]]
    binp <- p[state_name, !is.na(p[state_name, ])] #partial match

    while (nrow(binp) > 0) {

      state_name <-
        if (runif(1) < binp[1]) names(binp)[1] else names(binp)[2]

      binp <- p[state_name == rownames(p), !is.na(p[state_name, ])]
    }

    traj_s <- c(traj_s, state_name)
    traj_u <- c(traj_u, sutils[state_name])
  }
}

```

```
      cohort[[i]] <- traj_u
    }
  }
```

An example trajectory

```
cohort[[1]]
```

[illegible]

The mean summary statistics should be close to the true expected value. However, it appears to be pretty noisy and even for fairly large values (in terms of run time) it can be off by one or two.

```
mean(map(cohort, sum, na.rm = TRUE) %>% unlist())
```

```
## [1] 85.5559
```

C3. Markov-cycle tree

Given the following transition matrix

$$\begin{pmatrix} (1-\delta_0)\delta_c + \frac{1}{\delta_0 + (1-\delta_0)\delta_c} & \frac{0}{(1-\delta_0)(1-\delta_c)\beta_a} & \frac{0}{(1-\delta_0)(1-\delta_c)(1-\beta_a)} & \frac{0}{(1-\delta_0)(1-\beta_c)(1-\delta_a)} & \frac{0}{(1-\delta_0)(1-\beta_c)(1-\beta_a)} \\ \frac{\delta_0 + (1-\delta_0)\delta_c}{\delta_0} & \frac{(1-\delta_0)(1-\delta_c)\beta_a}{(1-\delta_0)\beta_c(1-\delta_a)} & \frac{0}{(1-\delta_0)\beta_c(1-\beta_a)} & \frac{0}{(1-\delta_0)(1-\beta_c)\beta_a} & \frac{0}{(1-\delta_0)(1-\beta_c)(1-\beta_a)} \end{pmatrix}$$

Then define the transition matrix object

```
p <- list()

p$dead <- c(1,0,0,0,0)

p$CA <-
  c(delta0 + (1-delta0)*deltac + (1-delta0)*(1-deltac)*deltaa, (1-delta0)*(1-deltac)*(1-deltaa),0,0,0)

p$cancer <-
  c(delta0 + (1-delta0)*deltac, (1-delta0)*(1-deltac)*betaa, (1-delta0)*(1-deltac)*(1-betaa),0,0)

p$AIDS <-
  c(delta0 + (1-delta0)*deltaa, (1-delta0)*betac*(1-deltaa), 0, (1-delta0)*(1-betac)*(1-deltaa), 0)

p$well <-
  c(delta0, (1-delta0)*betac*betaa, (1-delta0)*betac*(1-betaa), (1-delta0)*(1-betac)*betaa,
    (1-delta0)*(1-betac)*(1-betaa))

trans <- do.call(rbind, p)
```

Combine the tree data all together into a single list using a function.

```
create_tree <- function(trans, utils) {

  if (!all(rowSums(trans) == 1)) stop("probabilities don't sum to one")
  if (nrow(trans) != ncol(trans)) stop("not square matrix")
  if (nrow(trans) != length(utils)) stop("utils length doesnt match transition matrix dimensions")

  colnames(trans) <- rownames(trans)
  names(utils) <- rownames(trans)

  list(trans = trans,
        utils = utils)
}

my_tree <- create_tree(trans, sutils)
```

Check the input data.

```
str(my_tree)

## List of 2
## $ trans: num [1:5, 1:5] 1 0.10406 0.02615 0.08109 0.00118 ...
## .. attr(*, "dimnames")=List of 2
## .. ..$ : chr [1:5] "dead" "CA" "cancer" "AIDS" ...
## .. ..$ : chr [1:5] "dead" "CA" "cancer" "AIDS" ...
## $ utils: Named num [1:5] 0 0.3 0.6 0.5 1
```

```
##   ..- attr(*, "names")= chr [1:5] "dead" "CA" "cancer" "AIDS" ...
```

```
kable(my_tree$trans, digits = 3)
```

	dead	CA	cancer	AIDS	well
dead	1.000	0.000	0.000	0.000	0.000
CA	0.104	0.896	0.000	0.000	0.000
cancer	0.026	0.008	0.966	0.000	0.000
AIDS	0.081	0.002	0.000	0.916	0.000
well	0.001	0.000	0.003	0.008	0.988

Now we're ready to do the forward cycle. Basically, using the same approach as above for the separate probabilities, we simulate individuals and then take an average.

```
cohort <- list()
n_cohort <- 1000
p <- my_tree$trans

for (i in seq_len(n_cohort)) {

  traj_s <- NULL
  traj_u <- NULL
  state_name <- "well"

  while (state_name != "dead") {

    res <- rmultinom(n = 1, size = 1, prob = p[state_name, ])
    state_name <- rownames(res)[res[,1] == 1]
    traj_s <- c(traj_s, state_name)
    traj_u <- c(traj_u, sutils[state_name])
  }

  cohort[[i]] <- traj_u
}
```

Here's an example trajectory.

```
cohort[[1]]
```

```
## well well well well well well well well well well well well well well well well
## 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0
## well well well well well well well well well well well well well well well well
## 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0
## well well well well well well well well well well well well well well well well
## 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0
## well well well well well well well well well well well well well well well well
## 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0
## well well well well well well well well well well well well well well well well
## 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0
## well well well well well well well well well well well well well well well well
## 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0
## well well well well well well well well well well well well well well well well
## 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0
```

```
## well well well well well well well well well well well well well well well well
## 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0
## well well well well well well well well well well well well well well well well
## 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0
## well well well well well well well well well well well well well well well well
## 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0
## well well well well well well well well well well well well well well well well
## 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0
## well well well well well well well well well well well well well well well well
## 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0
## well well well well well well AIDS AIDS AIDS AIDS dead
## 1.0 1.0 1.0 1.0 1.0 0.5 0.5 0.5 0.5 0.0
```

The expected value is

```
mean(map(cohort, sum, na.rm = TRUE) %>% unlist())
```

```
## [1] 88.3194
```

C4. Regular Markov model

After initialising values for the calculation, for each cycle the probability of being in each of the states is calculated using the transition matrix and the previous cycle state occupancy probabilities. Similarly, the utilities associated with being in each state are calculated for each cycle.

```
run_model <- function(tree,
                        probs,
                        n_cycles = 1000) {

  if (!is.matrix(probs))
    probs <- matrix(probs, nrow = 1)

  qalys <- NULL
  costs <- NULL

  for (i in seq_len(n_cycles)) {

    probs <- rbind(probs, probs[i, ] %*% tree$trans)
    qalys <- rbind(qalys, probs[i, ]*tree$utils)
  }

  list(probs = probs,
       qalys = qalys)
}
```

By summing over all cycles we obtain the total utilities for each state. The total sum is the expected QALYs value for an individual starting in state `well` until `dead`.

```
init_pop <- c(0,0,0,0,1)
res <- run_model(my_tree, init_pop)
colSums(res$qalys)
```

```
##      dead      CA      cancer      AIDS      well
## 0.0000000 0.2134372 3.8587613 4.0724888 82.3270612
```

The total expected QALYs are therefore

```
sum(res$qalys)
```

```
## [1] 90.47175
```

C5. Roll back Markov-cycle tree

Let's write a recursive function to do the value iteration. Because this is more general than a simple binary tree we need to sum over the number of to-nodes at each step. Also, we need to limit the number of recursions since this function would run until we get a stack overflow error. In the original paper, Hazen (1992) gives a table for 1, 2, 3, 10, 100, 1000 cycles to show convergence. Here we include a `limit` argument which exits the function call after a certain tree depth is reached.

```
value_iteration <- function(n,      # starting node number
                           R,      # `reward` (utility/quality factor)
                           p,      # transition probability matrix
                           cycle,   # exits recursion after limit tree depth
                           limit = 100) {

  ## sub NAs so returns for absorbing state
  p[p == 0] <- NA
  p[p == 1] <- NA

  to_node <- which(!is.na(p[n, ]))

  if (length(to_node) == 0 || cycle == limit) {
    return(R[n])
  }
  else {
    Vsum <- 0

    for (i in seq_along(to_node)) {
      Vsum <- Vsum + p[n, to_node[i]]*value_iteration(to_node[i], R, p,
                                                       cycle = cycle + 1, limit = limit)
    }

    return(R[n] + Vsum)
  }
}
```

If we run this for the cycles in Table 2 in Hazen (1992), omitting the 1000 cycle because it takes too long to run, then we get the following same values

```
map_dbl(c(1,2,3,10,100),
        function(x) value_iteration(n = 5,
                                     R = my_tree$utils, p = my_tree$trans,
                                     cycle = 1, limit = x))
```

```
## [1] 1.000000 1.993599 2.980485 9.680872 63.018813
```

The problem with this representation is that it grows exponentially with the number of cycles. One way to speed things up is to do some of the calculation up-front so that we only do it once. We can achieve this by nesting a second function inside of the initial as follows.

```
value_iteration2 <- function(n,      # starting node number
                             R,      # `reward` (utility/quality factor)
                             p,      # transition probability matrix
```



```

        cycle,      # exits recursion after limit tree depth
        limit = 100) {

## sub NAs so returns for absorbing state
p[p == 0] <- NA
p[p == 1] <- NA

to_node <- map(seq(nrow(p)), function(i){ which(!is.na(p[i,])) } )

v_iter <- function(n, cycle) {

  if (length(to_node) == 0 || cycle == limit) {
    return(R[n])
  }
  else {
    Vsum <- 0
    for (i in seq_along(to_node[[n]])) {
      Vsum <- Vsum + p[n, to_node[[n]][i]]*v_iter(to_node[[n]][i], cycle = cycle + 1)
    }
    return(R[n] + Vsum)
  }
}

return(v_iter(n, cycle))
}

```

Although we still have the original main problem this does appreciably improve run time.

```

microbenchmark::microbenchmark(
  value_iteration(n = 5,
    R = my_tree$utils,
    p = my_tree$trans,
    cycle = 1, limit = 200),
  value_iteration2(n = 5,
    R = my_tree$utils,
    p = my_tree$trans,
    cycle = 1, limit = 200), times = 1)

```

Unit: seconds

##									expr
##	value_iteration(n = 5, R = my_tree\$utils, p = my_tree\$trans,								cycle = 1, limit = 200)
##	value_iteration2(n = 5, R = my_tree\$utils, p = my_tree\$trans,								cycle = 1, limit = 200)
##	min	lq	mean	median	uq	max	neval		
##	47.00826	47.00826	47.00826	47.00826	47.00826	47.00826	1		
##	18.71777	18.71777	18.71777	18.71777	18.71777	18.71777	1		